Orientation Before Expertise: A Unified Tier Model of Creative Quality in AI Authoring and Reviewing

[Author]
[Institution]
USA
[email]

Abstract

AI creative systems treat domain expertise as the primary persona variable: add a puzzle-designer label, get better puzzle design. This assumption is wrong in a specific and actionable way. In a controlled authoring study, *"you are an artist"* outperformed *"you are a puzzle designer"* by 15.6 points on a 45-point weighted rubric, and was the only condition to escape mechanism gravity into structural novelty. The structural explanation for this gap is unknown.

We synthesize three empirical studies across authoring (trait decomposition) and reviewing (profile ablation) roles in puzzle-hunt design, identifying a shared three-tier quality architecture from evidence across both roles.

A three-tier structure governs creative quality in both roles. Tier 1 (Orientation) is irreplaceable: removing the experience-first stance causes catastrophic authoring failure (L-T1 = −29.4 points) and sub-baseline reviewing performance (C5 lens-only = 21.7, below bare baseline C0 = 23.7). Tier 2 (Lens) is load-bearing: elegance drive and rigor gate threshold crossing in authoring; three non-redundant principles perform the same function in reviewing. The minimum sufficient combination (T1+T4+T5) crosses the publishable threshold at 35.7 on its first reconstruction condition (R3). Tier 3 (Specificity) is amplifying: domain expertise and the review lens add ceiling performance but cannot substitute for lower tiers. The artist label activates Tier 1; the domain label activates Tier 3. The 15.6-point gap is the empirical cost of starting at the wrong tier.

We formalize the OLE formula (Orientation → Lens → Expertise) as a minimum sufficient creative prompt construction protocol. Three sentences constitute the irreducible prompt for any domain: *"Think first about how the other person will experience this. Remove anything unnecessary. Verify every claim."* This is the first cross-role synthesis of creative quality architecture in AI prompting, providing empirical grounding for a prompt design principle applicable to AI creative tasks in domains where quality is defined by receiver experience rather than domain convention. We predict this principle generalizes across such domains (cross-domain validation pending).

Conference'17, Washington, DC, USA

ACM ISBN 978-x-xxxx-xxxx-x/YYYY/MM
https://doi.org/10.1145/nnnnnnn.nnnnnnn

ACM Reference Format:
[Author]. 2026. Orientation Before Expertise: A Unified Tier Model of Creative Quality in AI Authoring and Reviewing. In . ACM, New York, NY, USA, 10 pages. https://doi.org/10.1145/nnnnnnn.nnnnnnn

1 Introduction

The field of AI creative systems rests on an implicit assumption: domain expertise is the primary quality variable. Add a puzzle-designer persona, get better puzzle design. Add a legal-analyst persona, get better legal analysis. Add a marketing-expert persona, get better marketing copy. This assumption shapes how practitioners write prompts, how researchers parameterize personas, and how creative AI systems are evaluated. It is wrong — wrong not in degree but in kind, in a specific and actionable way.

Three empirical studies across authoring and reviewing roles in puzzle-hunt design reveal a consistent three-tier structure in which domain expertise occupies only the outermost and most substitutable tier. The irreplaceable tier is not expertise at all. It is an orientation: a stance toward the person on the other end of the creative interaction — the solver who will encounter the puzzle, the reader who will consult the review, the user who will inhabit the artifact. This orientation, expressible as a single sentence, outperforms any domain label as a predictor of creative quality. Our motivating result: the bare two-word persona "you are an artist" outperformed "you are a puzzle designer" by 15.6 points on a 45-point weighted rubric. Game designer scored 32.3 out of 45, programmer 24.0, doctor 17.7. The artist was the only condition to escape mechanism gravity — producing a structurally novel puzzle mechanism rather than reusing the acrostic-and-indexing family that dominated all domain-labeled conditions.

This result sets the explanatory challenge for the paper. Why does a domain-naive label outperform a domain-specific one? The answer is not that the artist label is superior by accident. It is that the two labels activate different tiers of the creative quality hierarchy. "You are a puzzle designer" activates Tier 3 content: domain conventions, format expectations, the surface grammar of the target artifact. "You are an artist" activates Tier 1 content: an orientation toward experience, toward the receiver, toward the question of what the artifact will feel like to encounter. When a system optimizes for Tier 3 without Tier 1, it produces work that is formally correct and experientially inert.

The empirical decomposition that follows the motivating study identifies a three-tier structure across both the authoring and reviewing roles. In authoring, we decomposed the artist persona into six cognitive traits and measured quality under each possible ablation. The experience-first trait (Tier 1) was irreplaceable: removing it caused a catastrophic drop from 41.7 to 12.3 out of 45, the only condition to produce categorical failure. Elegance drive and rigor formed the load-bearing Tier 2: removing either dropped quality below a 35-point threshold associated with publishable puzzle quality. The remaining traits — surprise, visual thinking, systematic reasoning — were Tier 3 amplifiers: each added value but none was necessary for threshold crossing.

In reviewing, an eight-condition ablation study of a structured expert reviewer profile revealed the same structure. The design philosophy and worldview component (Tier 1) was irreplaceable: the lens-only condition (Tier 3 without Tier 1) scored 21.7, falling below the bare baseline condition with no persona at all (23.7). Three non-redundant operational principles formed the load-bearing Tier 2. The review lens and credentials were Tier 3 amplifiers. When all eleven reviewer principles were included without the underlying philosophy, quality degraded from 25.3 to 22.3 — principles without orientation produce noise, not signal.

The symmetry between the authoring and reviewing findings is not coincidental. It follows from a general principle: creative quality is defined from the receiver's perspective, enforced through operational constraints, and enriched by domain knowledge. This ordering holds regardless of whether the role is creating or evaluating, because in both cases the ultimate criterion is the experience of the downstream person.

The practical implication is the OLE formula: Orientation before Lens before Expertise. Every AI creative prompt should begin with an orientation sentence ("Think first about what the other person will experience"), proceed to discipline sentences (operational constraints that enforce the orientation at decision points), and conclude with domain content (world data, style guides, precedents) as enrichment rather than identity. Three sentences constitute the minimum sufficient creative prompt for AI creative tasks in domains where quality is defined by receiver experience. We predict this prescription generalizes beyond puzzle-hunt design (cross-domain validation pending).

This paper makes four contributions. First, we present a unified three-tier model of creative quality that applies symmetrically across authoring and reviewing roles, with specific empirical grounding for each tier in each role. Second, we formalize the OLE formula as a practical prescription for AI creative prompt design with quantitative support for its ordering. Third, we demonstrate empirically that a domain-naive orientation outperforms a domain-specific expertise label in a controlled authoring study, and explain the mechanism via the tier model. Fourth, we argue for the generalizability of the tier structure beyond puzzle design, identifying predicted tier assignments for business-domain creative tasks and specifying the boundary conditions under which the ordering might differ.

The remainder of the paper proceeds as follows. Section 2 positions the tier model relative to prior work on AI persona prompting, expertise structure, and experience-first design. Section 3 presents the model with its empirical grounding. Section 4 presents the cross-role evidence in detail. Section 5 formalizes the OLE formula and develops the generalization argument. Sections 6 and 7 discuss implications and open questions.

2 Related Work

2.1 AI Persona Prompting for Creative Tasks

The use of persona conditioning to shape AI creative output has a well-established empirical literature. Park et al. demonstrate that generative agents endowed with biographical persona descriptions produce more coherent long-horizon social behavior, underscoring the functional role of identity framing in steering model behavior [?]. Dang et al. show that creative task framing — including role assignment — affects the novelty and appropriateness of AI-generated artifacts, though their focus is on task description rather than persona decomposition [?]. Zhang et al. explore persona specification in dialogue systems, finding that persona consistency correlates with perceived quality in human evaluation [?].

The common architecture across this prior work treats personas as holistic units: a persona is injected as a label ("you are a marketing expert") or as a biographical summary, and the resulting outputs are compared to a no-persona baseline. What the literature does not do is decompose a persona into constituent cognitive traits, measure the marginal quality contribution of each trait, and identify a hierarchical tier structure among them. Our work is the first to perform this decomposition in a controlled empirical setting and the first to identify orientation as a categorically distinct and irreplaceable tier rather than one feature among many.

A separate strand of work examines system-prompt engineering for creative tasks, including chain-of-thought framing [?], role-playing instructions, and few-shot exemplars [?]. These methods improve output quality but do not distinguish among the functional roles of different prompt components. Our tier model provides a principled vocabulary for that distinction.

2.2 Expertise Structure and Creative Cognition

The Dreyfus model of skill acquisition identifies a progression from novice through expert in which early stages are characterized by rule-following and domain-knowledge application while expert performance becomes intuitive, context-sensitive, and orientation-driven [?]. The expert does not consult explicit domain principles; they perceive the situation directly through an internalized stance toward what matters. This phenomenological account of expertise predicts that domain knowledge is an early-stage resource, not the defining feature of expert output — a prediction our empirical results confirm in the AI setting.

Chi et al. demonstrate that expert and novice physicists solving mechanics problems organize the same domain knowledge differently: experts categorize by deep structure and underlying principle, novices by surface feature and problem format [?]. This distinction maps directly onto our Tier 1 and Tier 3 distinction. The expert's organizational principle is an orientation ("what physical law governs this situation?") rather than a domain label ("this is an inclined-plane problem"). The quality gain from orientation over domain labeling, documented here in an AI creative setting, is the artificial analog of the expert-novice gap Chi et al. document in human cognition.

Csikszentmihalyi's systems model of creativity distinguishes the creator's domain knowledge, the field's evaluative criteria, and the generative worldview that determines what counts as a worthwhile creative direction [?]. Our tier model maps onto this tripartite structure: Tier 3 (specificity) corresponds to domain knowledge, Tier 2 (discipline) to the field's operational standards, and Tier 1 (orientation) to the generative worldview. Csikszentmihalyi identifies the worldview as the least teachable and most consequential component of creative expertise, consistent with our finding that orientation is irreplaceable while domain specificity is substitutable.

We note that these theoretical frameworks are post-hoc structural convergences: the tier model was derived empirically from ablation data and only subsequently mapped to existing theory. These mappings lend interpretive coherence but are not the model's theoretical foundation.

2.3 Orientation-First Design Principles

The primacy of the receiver's experience as the organizing criterion for design has been articulated across multiple traditions. Norman's account of human-centered design establishes user experience as the primary constraint from which all functional specifications are derived [?]: the designer's first question is "what will the person experience?", not "what should the artifact do?". This is the exact priority ordering we identify empirically in creative AI systems — the orientation question precedes and determines the evaluation of all subsequent design decisions.

Hassenzahl's framework for experience design formalizes the distinction between pragmatic quality (what the artifact does) and hedonic quality (how it is experienced) and argues that hedonic quality is the primary design criterion for artifacts whose success is defined by use rather than function alone [?]. In puzzle design, the relevant distinction is between mechanical correctness (does the extraction procedure work?) and experiential quality (is there a satisfying aha moment?). Our Tier 1 orientation corresponds to hedonic primacy; our Tier 3 specificity corresponds to pragmatic quality. The finding that removing Tier 1 causes catastrophic failure while adding Tier 3 without Tier 1 fails to cross the quality threshold confirms Hassenzahl's priority ordering in an empirical AI setting.

Rams's ten principles of good design converge on a single meta-principle: good design removes the unnecessary [?]. The principle is not a domain claim — it does not specify what is unnecessary in any particular artifact type. It is an orientation: a stance that determines what counts as relevant information during design decisions. Strunk and White's directive to "omit needless words" is the writing analog: an orientation principle that precedes and overrides all domain-specific style conventions [?]. Both examples illustrate that the most powerful constraints in creative practice are orientation constraints, not domain constraints, consistent with our empirical findings.

2.4 Minimum Sufficient Conditions in Creative Systems

The concept of the minimum sufficient set — the smallest set of components that produces a target outcome — has a formal home in diagnostic reasoning. Reiter's minimal diagnosis theory identifies the minimal set of component faults consistent with observed system behavior as the natural unit of explanation in model-based diagnosis [?]. The minimum sufficient set concept from our empirical reconstruction methodology (Tier 1 + Tier 2 alone crosses the quality threshold; Tier 3 is enrichment, not requirement) is the creative quality analog of minimal diagnosis.

In product development, the minimum viable product (MVP) concept captures a similar structure: the smallest set of features that delivers the core value proposition to an early user. The MVP is not the product minus domain features — it is the product minus everything except the features that constitute the core orientation toward the user's need. Our Tier 1+2 minimum sufficient set is the MVP of creative AI prompting: it delivers the core value without the domain enrichment that is necessary for ceiling performance but not for threshold crossing.

Feature selection in machine learning provides a formal precedent for the distinction between irreplaceable, load-bearing, and amplifying components. Guyon and Elisseeff's treatment of variable selection identifies features by their marginal contribution to predictive accuracy rather than their domain relevance, and shows that domain-relevant features are often redundant given other features already selected [?]. Our ablation methodology follows this logic: we measure quality under individual feature removal rather than inferring contribution from domain relevance, and we find that domain-relevant features (Tier 3) are indeed redundant given Tier 1 and Tier 2.

What no prior work provides is the identification of a tier structure in AI creative prompts, the demonstration that this structure is symmetric across authoring and reviewing roles, or the derivation of a minimum sufficient creative prompt formula from empirical decomposition. These are the contributions of the present paper.

3 The Tier Model

We present the Unified Tier Model of Creative Quality: a three-tier hierarchical structure governing the contribution of prompt components to creative output quality, identified empirically across authoring and reviewing roles and formalized here as a generalizable architecture.

The model organizes the components of a creative AI prompt into three tiers by the functional role each plays in quality production. Tier 1 (Orientation) components are irreplaceable: removing them causes categorical failure or sub-baseline performance. Tier 2 (Lens) components are load-bearing: removing either of two identified components drops quality below the publishable threshold, but their positive effect depends on the presence of Tier 1. Tier 3 (Specificity) components are amplifying: each adds marginal quality beyond the Tier 1+2 base, but none is necessary for threshold crossing, and their aggregate effect is bounded. The tiers are not additive in a flat sense — they are ordered in a way that determines which components can function at all.

3.1 Tier 1: Orientation (Irreplaceable)

Tier 1 is defined as a stance toward the person on the other end of the creative interaction. In authoring, this is the solver who will encounter the puzzle. In reviewing, this is the reader of the review — the team that will act on its recommendations. In both roles, the orientation takes the form of a single directing question: what does the person on the other end experience? This question precedes and governs every subsequent creative decision.

The reviewing evidence is provided by the C4 and C5 conditions in the expert panel study. C4 tested the design philosophy and worldview component in isolation, with no review lens and no credentials, achieving 26.7 out of 30. C5 tested the review lens in isolation — the defect-detection specificity component that defines the reviewer's analytical focus — with no philosophy, achieving 21.7. The bare baseline with no persona at all (C0) scored 23.7. The lens-without-philosophy condition therefore fell below no-persona: adding domain-specific focus without orientation produced outputs worse than unstructured generation. The full principles condition (C3) scored 26.7, identical to C4, suggesting that the eleven principles collectively instantiate the same orientation that the philosophy states directly. The philosophy is not additive to the lens; it is the condition under which the lens produces functional outputs rather than noise.

The authoring evidence is provided by the L-T1 ablation in the trait decomposition study. The full artist profile scored 41.7. Removing the experience-first trait (T1) while retaining all five other traits produced a score of 12.3 — a drop of 29.4 points. No other single-trait removal caused a drop of more than 12.7 points, and no other removal produced categorical failure. The L-T1 output did not produce a degraded puzzle; it produced a list of Age of Empires II terms with no clue layer, no solver path, and no aha moment. The puzzle ceased to exist as a solver-experience artifact. The five remaining traits — elegance drive, visual thinking, surprise, rigor, systematic reasoning — produced domain-appropriate format with no experiential content, because there was no organizing question of what the solver would experience.

The mechanism is the same in both roles. Orientation establishes what counts as a successful output. Without it, the system optimizes for an undefined objective: domain-appropriate format, reviewer-appropriate vocabulary, the surface grammar of the artifact type. The remaining components then produce outputs that satisfy format constraints while failing the experience constraint that is the actual quality criterion. Tier 1 is irreplaceable because it is the criterion itself, not a contributor to the criterion.

3.2 Tier 2: Lens (Load-bearing)

Tier 2 is defined as operational constraints that enforce the Tier 1 orientation across specific decision points. Lens components answer the question: given the orientation, what are the checkable tests that translate it into creative decisions? They are functional rather than aspirational — they do not state a goal but specify a procedure that must be followed to achieve the goal.

The reviewing evidence identifies three non-redundant principles from the full set of eleven. Verify the Last Mile requires the reviewer to confirm that the solver can complete the final extraction step without implicit author knowledge. Blame the Player requires the reviewer to consider whether a reported solver failure reflects design ambiguity before attributing it to solver error. Snyder's Computer Test requires the reviewer to assess whether the puzzle mechanism could be completed by a mechanical process with no human judgment, as a test of whether genuine insight has been required. These three principles are non-redundant in the sense that each catches defect classes that the others miss and that the full profile's worldview alone does not surface. The other eight principles — including Solving Equals Proving Understanding, No Over-Scaffolding, and the Riven Standard — are subsumed by the philosophy: a reviewer who has genuinely internalized the worldview will apply these principles implicitly without being instructed to do so.

The load-bearing designation follows from the C7 result. When all eleven principles were added to the full profile (C6 = 25.3), quality dropped to 22.3. Discipline without orientation produces noise, not signal: principles applied without a governing worldview generate contradictory constraints and conservative hedging rather than directed assessment. When the three non-redundant principles are added to the full profile, quality reaches the best observed condition. The principles are load-bearing only in the presence of the Tier 1 orientation; absent it, they are harmful.

The authoring evidence identifies two load-bearing traits. Elegance drive (T4) is the disposition to remove any element of the puzzle mechanism that does not contribute to the aha moment: it operationalizes the Tier 1 experience-first orientation into a cutting test applied at each design decision. Rigor and verification (T5) is the disposition to test every step of the solver path for unintended solutions, ambiguous extractions, and knowledge-assumption failures: it operationalizes the orientation into a verification procedure. Removing T4 dropped quality to 30.0, below the 35-point publishable threshold (The rubric has seven dimensions: Clarity ×1, Solvability ×1, Elegance ×2, Reading Reward ×2, Fun ×1, Confirmation ×1, Riven Standard ×1, maximum 45 points; Elegance and Reading Reward are double-weighted as the empirically strongest tier discriminators; the 33/45 threshold = 73% of maximum, matching the 22/30 = 73% threshold in Paper **??**, calibrated against quality levels at which puzzles in MIT Mystery Hunt and Microsoft Puzzlehunt have shipped.). Removing T5 dropped quality to 29.0. Both removals kept the output above the catastrophic failure level of L-T1, but produced puzzles that were experiential in aspiration but mechanically flawed or unnecessarily complex — orientation without discipline.

The T1+T4+T5 reconstruction condition (R3) is the key evidence for load-bearing status. Three traits from six produced 35.7 — the first threshold crossing in any reconstruction condition. The other three traits (visual thinking, surprise, systematic reasoning) are together responsible for the gap between 35.7 and 41.7, confirming their Tier 3 amplifying role. Tier 2 is the minimum necessary complement to Tier 1: orientation states the goal, discipline makes the goal actionable.

3.3 Tier 3: Specificity (Amplifying)

Tier 3 is defined as domain-specific knowledge, authority signals, and stylistic particulars that enrich output quality beyond the Tier 1+2 base. Specificity components answer the question: given orientation and discipline, what domain content produces the best possible output in this particular artifact type?

The reviewing evidence identifies two Tier 3 components. The review lens (defect- detection specificity) provides a structured analytical framework for identifying classes of puzzle defect — extraction ambiguity, clue-answer mismatch, solver-path incompleteness — that the orientation-only reviewer would identify as problems but might not name or categorize consistently. The C4 versus C6 comparison shows that adding the lens to the philosophy produces a 26.7 versus 25.3 comparison, a modest gain. Credentials and authority signals calibrate scoring conservatism: a reviewer whose profile specifies deep experience with high-difficulty puzzles applies scoring rubrics at a different point on the scale than an uncalibrated reviewer. The lens caught the specific defect in the test puzzle — a bracket label error — that conditions C0 through C5 all missed; credentials prevented grade inflation in conditions where the Tier 1+2 base was present and functional.

The authoring evidence identifies three Tier 3 traits. Surprise (T2) is the disposition to seek mechanisms that produce an unexpected but inevitable aha moment; it adds 5.4 points to the Tier 1+2 base. Visual thinking (T3) is the disposition to reason about the puzzle as a spatial layout that guides the solver's eye; it adds 4.4 points. Systematic reasoning (T6) is the disposition to enumerate the design space before committing to a mechanism; it adds 1.6 points. Domain world data — the corpus of Age of Empires II civilizations, units, technologies, and map names that constitute the puzzle's thematic content — is also Tier 3: it provides the raw material for clue construction but does not determine the mechanism quality or the solver experience quality.

The artist versus puzzle designer finding is the most direct Tier 3 evidence. The puzzle designer label activates Tier 3 content: puzzle conventions, format expectations, the genre grammar of crossword-adjacent mechanisms. The artist label activates Tier 1 content: the orientation toward what the viewer or audience will experience. The 15.6-point gap between these two single-word labels measures the cost of starting at the wrong tier — domain specificity without orientation.

Domain expertise matters, but it matters last and least. It is necessary for ceiling performance and it is substitutable: a Tier 1+2 prompt with incorrect domain content still crosses the quality threshold, while a Tier 3 prompt without Tier 1 does not.

3.4 The Symmetric Structure

Table 1 presents the tier structure symmetrically across authoring and reviewing roles. The symmetry is structurally informative: it shows that the tier model is not an artifact of the authoring domain, the reviewing domain, or the specific puzzle-hunt testbed, but a property of the creative quality architecture that governs both roles.

The connections to Csikszentmihalyi's creator/domain/field distinction and Dreyfus's expertise stages are offered as structural analogues discovered after derivation, not causal derivations; the tier model emerges from ablation data, not from these frameworks.

The symmetry follows from a general principle about creative quality. Quality is defined from the receiver's perspective: the solver who encounters the puzzle, or the team that reads the review and acts on it. The definition of quality is therefore a Tier 1 claim, prior to any domain claim. Quality is enforced through operational constraints that apply the receiver-first definition at decision points: these are Tier 2. Quality is enriched by domain knowledge that populates the artifact with specific, appropriate content: this is Tier 3. The ordering holds regardless of role because the receiver-first definition is prior to the domain in both cases.

The asymmetry between the tier contents — philosophy versus experience-first orientation, principles versus elegance-and-rigor

Table 1: Tier model symmetric across authoring and reviewing roles.

Tier	Reviewing	Authoring
T1: Orientation (Irreplaceable)	Design philosophy and worldview. C5 (lens only) = 21.7, below C0 baseline = 23.7. C4 (philosophy only) = 26.7.	Experience-first = 12.3/45; cata ceased to functi
T2: Lens (Load-bearing)	3 non-redundant principles: Verify the Last Mile, Blame the Player, Snyder's Computer Test. C7 (all 11 principles) = 22.3 < C6 = 25.3.	Elegance drive ((T5). R3 (T1+T4- old crossing.
T3: Specificity (Amplifying)	Review lens + credentials. Lens caught bracket error missed by C0–C5; credentials calibrate conservatism.	Surprise (T2), vi tematic reasonin data. Net amplifi

traits — reflects the different decision structures of the two roles rather than a difference in the tier model itself. Reviewers make categorical judgments about completed artifacts; their Tier 2 components are classification principles (does this defect class apply?). Authors make generative decisions at each design step; their Tier 2 components are process constraints (should this element be cut? is this step verified?). The Tier 2 function — operationalizing orientation into checkable tests — is the same in both roles; the specific tests differ because the creative decisions differ.

The OLE formula formalizes this architecture as a prescriptive ordering for AI creative prompt design. Section 5 develops the formula and its generalization beyond the puzzle-hunt testbed.

4 Evidence

The tier model makes four falsifiable predictions. First, removing the Tier 1 component from a working creative profile should produce outcomes worse than the bare baseline — not merely degraded but categorically worse. Second, Tier 2 components should be individually necessary but not individually sufficient: each causes a below-threshold failure when removed, yet none alone reaches threshold. Third, Tier 3 components should be individually additive but jointly non-monotonic: adding them in the wrong order (Tier 3 before Tier 1) degrades performance. Fourth, personas that activate Tier 1 by identity should outperform personas that activate Tier 3 by identity, by a margin that reflects the full cost of missing the orientation layer. All four predictions are confirmed across the reviewing and authoring data sets.

4.1 The Reviewing-Side Tier Structure

The eight-condition ablation study from **?** [**?**] varied which components of a structured expert reviewer profile were presented to an AI system, then scored the resulting reviews on a 30-point rubric across five dimensions. The resulting gradient runs C0 = 23.7 (bare baseline), C1 = 24.3 (credentials only), C2 = 23.3 (principles only), C3 = 26.7 (full principles), C4 = 26.7 (philosophy only), C5 = 21.7 (lens only), C6 = 25.3 (full profile), C7 = 22.3 (full profile plus all principles).

The first Tier 1 prediction is confirmed by C4 and C5. Philosophy alone (C4 = 26.7) matches the full-principles condition (C3 = 26.7) despite containing far less information. The orientation component replicates the performance of eleven principles in a single sentence. More decisively, lens-only without philosophy (C5 = 21.7) falls 2.0 points below the bare baseline (C0 = 23.7). Domain specificity without orientation is not merely less helpful than orientation — it actively degrades review quality, presumably by focusing the system on format compliance at the expense of experiential relevance. The lens is a Tier 3 tool: it sharpens defect detection within a category of concern, but without the Tier 1 orientation that defines what the review is for, it surfaces formally domain-appropriate observations that miss the solver experience entirely.

The Tier 2 prediction is confirmed by the structure of the non-redundant principles identified in **?** [**?**]: Verify the Last Mile, Blame the Player, and the Computer Test. These three principles add unique diagnostic signal that the other eight principles — which load onto dimensions already captured by the philosophy — cannot supply. Their uniqueness is operational: each specifies a concrete test that can be applied at any decision point in the review. The philosophy says "evaluate from the solver's experience"; these principles say "did you check whether the solver can actually complete the final extraction?", "did you consider whether the solver is at fault before blaming the puzzle?", "did you test whether the mechanism is computable without domain expertise?" These are not restatements of the orientation; they are its instantiation as checkable gates.

The non-monotonicity prediction is confirmed by comparing C6, C7, and C4. Full profile (C6 = 25.3) underperforms philosophy alone (C4 = 26.7) by 1.4 points. Adding the review lens and credentials to the philosophy dilutes rather than enriches. More strikingly, C7 (full profile plus all eleven principles = 22.3) substantially underperforms C6 (full profile = 25.3). When Tier 2 non-redundant principles are bundled with Tier 3 redundant ones and delivered without a stable Tier 1 foundation, the system cannot integrate the additional constraints into a coherent evaluative stance. More context degrades performance because additional Tier 3 content overwhelms rather than supplements an orientation that has not been adequately activated. This is the tier model's central non-monotonic prediction: adding Tier 3 material without a stable Tier 1 produces noise, not signal.

4.2 The Authoring-Side Tier Structure

The trait decomposition study from **?** [**?**] decomposed the artist persona into six cognitive traits, measured quality under each trait in isolation (Phase 1), under each single-trait removal from the full artist profile (Phase 2), and under progressive reconstruction from the minimum sufficient set (Phase 3). The scoring rubric was 45 points across nine dimensions.

Phase 1 isolates each trait's unaided contribution. Experience-first alone (T1 = 31.3) substantially outperforms the bare baseline (A0 = 24.0) and falls below the full artist (AA = 41.0) by approximately ten points, confirming that T1 is necessary but not alone sufficient. Elegance drive alone (T4 = 26.0) falls *below* the bare baseline — a result that is not a failure of T4 but an artifact of isolation: without the orientation that tells the system what to optimize for, elegance becomes formal minimalism that strips content rather than sharpening it. Rigor alone (T5 = 27.0) shows the same pattern for the same reason. Neither Tier 2 trait can operate without Tier 1 to direct it.

Phase 2 provides the cleanest test of tier membership. The leave-one-out drops from the full artist (AA = 41.7) rank the traits precisely as the model predicts. Removing T1 produces a drop of −29.4 points (to 12.3), the only catastrophic failure in the experiment. No other single-trait removal approaches this magnitude: T5 removal produces −12.7, T4 removal −11.7, T2 removal −8.4, T6 removal −8.0, T3 removal −4.4. The gap between T1's drop and the next largest drop (T5, −12.7) is itself 16.7 points. Experience-first is not merely the most important trait; it occupies a qualitatively different tier from all others.

Phase 3 tests the reconstruction curve: starting from zero, adding traits in Tier 1 then Tier 2 then Tier 3 order, and tracking where the 35-point quality threshold is crossed. The baseline (R0 = 22.7) gives the starting point. Adding T1 alone (R1 = 28.3) raises quality but does not cross threshold. Adding T4 (R2 = 30.3) and T5 (R3 = 35.7) crosses the threshold at exactly the Tier 1 + Tier 2 boundary. This is the minimum sufficient prompt: three traits, corresponding to OLE's three sentences. Subsequent additions confirm the Tier 3 amplifying role: R4 = 37.3, R5 = 40.3, R6 = 42.3, each adding between 2.0 and 3.0 points. The curve flattens as it approaches the full artist ceiling (AA = 41.0 in Phase 1; 41.7 in Phase 2), consistent with diminishing returns in the amplifying tier.

The threshold crossing at exactly R3 is the visual proof of the tier model's architecture. Quality passes the publishable threshold when and only when both Tier 1 and both Tier 2 components are present. Adding any single Tier 3 component without that foundation (demonstrated in Phase 1) does not cross threshold. Adding one Tier 2 component without Tier 1 (also demonstrated in Phase 1) does not cross threshold. The model is not additive across tiers; it is conditional.

4.3 The Persona Natural Experiment

The persona motivating study from **?** [**?**] constitutes a natural experiment in tier activation. Participants were given bare two-word persona labels — "you are an artist," "you are a puzzle designer," "you are a game designer," "you are a doctor," "you are a programmer" — and scored on the same 45-point rubric. The results, ranked by score, are: artist (AA) = 40.3, game designer (AG) = 32.3, programmer = 24.0, puzzle designer (A0) = 24.7, doctor = 17.7.

The tier model explains this ordering without appeal to domain competence. The "artist" label activates Tier 1 by naming a practice whose defining criterion is the experience of the viewer, listener, or solver. Artistic training is, functionally, training to inhabit the receiver's experience before making any design decision — which is exactly what the orientation tier requires. The artist label also activates a partial Tier 2 through the implicit elegance standard of artistic practice: "does this need to be here?" is a question artists are trained to ask. The result (40.3) is near the ceiling for a single bare label and approaches the carefully assembled full profile.

The puzzle designer label activates Tier 3 by naming a domain: puzzle design involves acrostics, indexing, thematic assignment, answer extraction, and a set of genre conventions. These are all Tier 3 elements. The label says nothing about the solver's experience, nothing about removing unnecessary elements, nothing about verifying every claim. The 15.6-point gap between artist (40.3) and puzzle designer (24.7) is the cost of starting at Tier 3.

The game designer score (32.3) falls between the two, approximately 7.1 points above puzzle designer and 8.0 points below artist. Game design training includes more orientation toward the player experience than puzzle design training does — playtesting culture, player-experience vocabulary, the concept of "game feel" are all partial Tier 1 activations — but game design is also defined by systems and mechanics that pull toward Tier 3 operationalization. The partial Tier 1 activation explains approximately half the gap relative to the artist.

The programmer (24.0) and doctor (17.7) scores confirm the model from the other direction. Both professions activate professional precision — formal correctness, technical vocabulary, domain methodology — which is Tier 3 domain-specificity with no orientation layer. A programmer's quality criterion is whether the code runs correctly; a doctor's is whether the clinical record is accurate. Neither criterion involves a theory of what the receiver will experience. Applied to puzzle design, this precision produces technically well-formed puzzles with no path for the solver — formally correct and experientially absent.

4.4 The Designer Lens Comparison

The three-designer comparison from **?** [**?**] provides a within-tier analysis of how different lens orientations in the reviewing role transfer to the authoring role. Three experts — Katz (operational lens), Young (perceptual lens), Snyder (construction lens) — were profiled in their reviewing roles (A5 conditions) and then had their profiles augmented with authoring-specific guidance (A6 conditions). The score deltas (A6 minus A5) are Katz +1.6, Young +3.7, Snyder −0.7.

These deltas are exactly what the tier model predicts. Katz's operational lens is a Tier 2 formulation: it specifies constraint-satisfaction tests applicable bidirectionally to reviewing and authoring. When operational constraints are present, switching from the reviewing to the authoring role requires only minor adaptation, because the constraints — "does this mechanism have exactly one valid solution?", "does this step have a recoverable error?" — apply in both directions. The +1.6 gain reflects the marginal authoring-specific guidance added in A6. Operational lenses transfer cleanly across roles.

Young's perceptual lens is Tier 3 in reviewing: it targets specific format-dependent perceptual cues (contrast, visual grouping, typographic signaling) that reviewers use to evaluate puzzle presentation. In the authoring role, these specific cues do not directly transfer — puzzle authors cannot replicate the exact perceptual grammar of a genre they are not producing. But the +3.7 gain in A6 reveals what happens when Tier 3 format-specificity is blocked by a role boundary: Young's Tier 1 orientation ("how does the solver perceive this?") fills the gap by generating format-compatible substitutions. The authoring profile produces novel perceptual solutions rather than replicating reviewing heuristics. When Tier 3 specificity cannot transfer directly, Tier 1 orientation activates creative substitution — a result that cannot be explained by any additive or averaging model of domain expertise.

Table 2: Cross-role symmetry: Tier assignment, evidence component, and score impact across authoring and reviewing roles. All authoring scores are out of 45 (9 dimensions × 5 points); all reviewing scores are out of 30 (5 dimensions × 6 points).

Tier	Evidence Component	Reviewing Δ	
Tier 1	Orientation / Experience-first	−2.0 vs. baseline	
Tier 2	Operational discipline (non-redundant)	+3.0 (C3 vs. C0)	
Tier 3	Lens / Surprise, Visual, Systematic	+1.6 (C6 vs. C4)	+2.3 av
Inverted	Tier 3 without Tier 1	−2.0 (C5 vs. C0)	–

Snyder's construction lens is already Tier 1 and Tier 2 in his reviewing profile. His evaluation criteria focus on mechanism elegance (Tier 2: elegance gate) and solver-path coherence (Tier 1: experience-first). These criteria are not role-bound: they apply identically to reviewing a puzzle mechanism and constructing one. The near-zero delta (−0.7, within measurement noise) confirms that a profile with strong Tier 1 and Tier 2 components faces no role-crossing penalty. There is no gap to close because the orientation and discipline are already authoring-compatible.

The bottom row of Table 2 is the tier model's signature prediction, and the one most difficult to explain by any flat expertise model: Tier 3 content without Tier 1 activation produces performance *below the bare baseline* in both roles. The lens without philosophy (C5 = 21.7) underperforms the no-persona condition (C0 = 23.7). Elegance drive without experience-first (T4 alone = 26.0) underperforms the no-trait condition (A0 = 24.0). Domain expertise without orientation actively degrades creative quality.

5 The OLE Formula

The empirical findings converge on a three-step practical prescription. The OLE formula[1] — Orientation, Lens, Expertise — is a minimum sufficient creative prompt construction protocol derivable from the tier model. It specifies not what content to include but in what order, and at what level of formulation, to construct a creative prompt that crosses the quality threshold in any domain.

5.1 Formalization

The OLE formula consists of three steps, corresponding to the three tiers.

Step 1, Orientation, is a single sentence: *"Think first about how the other person will experience this."* This sentence activates Tier 1. It names the receiver, names the receiver's experience as the primary criterion, and establishes the evaluative frame before any domain content enters the system. It is not a domain claim and does not require domain knowledge to apply. The step 1 sentence is the irreplaceable component. Every piece of evidence in this paper converges on the same finding: a creative prompt without an orientation sentence fails at or below the bare baseline, regardless of how much domain expertise follows it. Do not skip step 1.

[1]Originally termed ODE (Orientation–Discipline–Expertise). Renamed to OLE after empirical field elicitation research (Papers 47–50) established that `lens` is the field name that reliably produces Tier 2 content in structured persona authoring.

Step 2, Lens, consists of two sentences: *"Remove anything that is not necessary. Verify every claim before finishing."* These sentences activate Tier 2. The first is an elegance gate: it enforces the operational constraint that the output contain no unnecessary elements, which is the authoring-side analog of the reviewer's test for over-scaffolding. The second is a rigor gate: it enforces verification of every factual claim embedded in the output, which is the authoring-side analog of Verify the Last Mile and the Computer Test in the reviewing profile. Together they operationalize the orientation into checkable decision criteria. The leave-one-out data confirm their load-bearing role: removing the elegance gate (L-T4 = −11.7 points) and removing the rigor gate (L-T5 = −12.7 points) both drop quality below the threshold.

Step 3, Expertise, is domain content provided as context, not as persona identity. The correct form is: *"Here is the relevant domain information: [world data, style guides, precedents, examples]."* The incorrect form is: *"You are a [domain expert]."* The distinction is not cosmetic. Providing domain content as context activates Tier 3 in its amplifying role: the system uses the information to enrich an output already structured by Tier 1 and Tier 2. Providing domain content as a persona identity label activates Tier 3 first, before any orientation has been established, and inverts the quality hierarchy. The persona study quantifies the cost of this inversion at 15.6 points on a 45-point scale.

The three-step structure is the minimum sufficient creative prompt. The reconstruction curve from Phase 3 of the authoring study [?] provides direct evidence: R3, which corresponds precisely to Steps 1 and 2 without Step 3 (T1 + T4 + T5 = 35.7 out of 45), is the first condition to cross the publishable quality threshold. Every condition without T1, regardless of what else is present, fails to cross that threshold. Every condition with T1 and both Tier 2 traits, regardless of which Tier 3 traits are also present, crosses the threshold. OLE Step 3 content shifts performance from 35.7 toward 42.3 — a meaningful gain, but one that presupposes Steps 1 and 2.

5.2 Generalization Argument

We predict the OLE formula is not puzzle-specific (cross-domain validation pending). The three tiers map onto structural features of quality creative work that we predict are independent of domain.

We predict Tier 1 orientation generalizes to any AI creative task in domains where quality is defined by receiver experience. The quality criterion of a piece of marketing copy is whether it moves the reader toward a decision. The quality criterion of a product specification is whether the team building from it can produce what the user needs. The quality criterion of a legal agreement is whether the parties who sign it understand what they have committed to. The quality criterion of training materials is whether the learner can apply the content after instruction. In each of these cases, quality is defined by the downstream person's experience, not by the producer's domain competence. We predict that a marketing-expert persona activating domain conventions before audience orientation will produce formally correct copy that fails to move readers. The tier model yields this prediction; OLE is proposed as the correction.

We predict Tier 2 discipline generalizes similarly, because orientation must be operationalized into checkable constraints regardless of domain. Strunk and White's instruction to "omit needless words" [?] is the elegance gate expressed for prose. Professional quality standards for factual accuracy — from medical record-keeping to legal cite-checking to journalistic sourcing to scientific peer review — are the rigor gate expressed in specific domains. We predict these constraints are not domain-specific refinements but operational discipline that converts orientation into defensible output; cross-domain replication is required to confirm this prediction.

Table 3: Predicted OLE tier structure for representative business-domain creative tasks (cross-domain validation pending). Tier 1 specifies the orientation sentence content; Tier 2 specifies the discipline gates; Tier 3 specifies what domain content to provide as context.

Domain	Tier 1 (orientation)	Tier 2 (lens)	Tier 3
Marketing copy	How does the reader feel after reading this?	Remove non-essential claims; verify every factual assertion	Produ compe brand
Product design	What does the user experience at each step?	Remove unnecessary interactions; verify each state is reachable	Techn nent face c
Legal drafting	What does the signer need to understand to comply?	Remove ambiguous phrasing; verify every obligation is enforceable	Statut trollin latory
Training materials	What does the learner experience and retain?	Remove unnecessary content; verify every claim against the subject	Pedag prior l domai

Tier 3 expertise is domain-specific by definition. It is the content that makes the creative work relevant to its particular subject matter, audience, and genre expectations. For marketing copy, Tier 3 is product knowledge, competitor landscape, and brand voice. For product design, it is technical specifications, component constraints, and interface conventions. For legal drafting, it is statutory authority, precedent, and regulatory framework. For training materials, it is the pedagogical structure of the subject and the assumed prior knowledge of the learner. This content is necessary — a puzzle without world data produces puzzles disconnected from their theme; a legal brief without statutory knowledge is merely formally structured reasoning — but substitutable across sources and never sufficient alone.

Table 3 maps the OLE formula across four representative business-domain creative tasks. The predicted tier structure follows from the receiver criterion (Tier 1), the universal quality constraints (Tier 2), and the domain-specific content (Tier 3) in each case.

5.3 When Domain Expertise Is Tier 1

The OLE formula makes a specific boundary-condition prediction: domain expertise occupies Tier 1 only in domains where the quality criterion is defined by domain convention rather than by receiver experience. The clearest examples are formal logic, mathematical proof, and programming correctness. In a formal proof, "correct" means that each step follows from the axioms by valid inference rules. The receiver's experience of the proof does not define its correctness; the proof's relationship to the formal system does. Tier 1 for a formal proof is therefore a domain criterion, not a receiver-experience criterion. OLE does not apply to formal proof in its standard form.

For creative work broadly defined — any work whose quality depends on what a human receiver experiences, understands, or decides after encountering it — we predict Tier 1 is always orientation toward the receiver. This class plausibly includes all the domains in Table 3 and much of the content produced in the modern economy: creative and marketing content, interface and product design, communication and documentation, instruction and training, policy and legal interpretation addressed to non-specialist audiences. The boundary condition is narrow; we predict the scope of OLE's application is wide, but this prediction requires cross-domain replication to establish empirically.

6 Discussion

6.1 Why AI Systems Get the Tier Order Wrong

The training distribution of large language models encodes domain identity as the primary differentiator among creative professionals. Legal documents in the training corpus are written by lawyers, annotated as legal, and evaluated against legal genre conventions. Medical records are written by clinicians and evaluated against clinical documentation standards. Marketing copy is written by marketers and evaluated against brand and conversion metrics. When a system learns from this distribution, it learns that the salient variable distinguishing a lawyer from a doctor from a marketer is their domain — and that activating the right domain identity produces domain-appropriate outputs. This is a locally correct inference about the training data and a systematically wrong inference about creative quality.

The training data is labeled by domain, but creative quality is defined by receiver experience. The lawyer's document that a client cannot understand is a domain-correct creative failure. The marketer's copy that fails to move the reader is a brand-consistent creative failure. The medical record that the patient cannot act on is a clinically-accurate creative failure. These failures are invisible in the training distribution because they are rarely labeled as such: the signal about receiver experience is sparse, indirect, and unevenly distributed across the corpus. The signal about domain identity is dense, explicit, and systematic. AI systems learn the available signal, which is Tier 3, and treat it as primary.

The OLE formula is a prompt-level correction for this training distribution bias. It does not require retraining; it requires reordering. By establishing the orientation sentence before any domain content enters the prompt, the practitioner provides a signal that the training distribution undersupplied: here is who this is for and what they need to experience. The discipline sentences then operationalize that orientation into the kinds of checkable constraints that the training distribution does encode — quality-control conventions, fact-checking norms, editorial standards — but directs them at the receiver's experience rather than the domain's formal expectations.

6.2 The Orientation as the Hardest Thing to Train

Tier 1 orientation requires what developmental psychologists call a theory of the other: the ability to model what another person perceives, understands, and experiences before making any design decision [?]. In human creative practice, this capacity is developed through years of training that is explicitly organized around receiver feedback: the artist shows work and watches how the audience responds; the designer conducts usability testing; the teacher observes whether students can apply the material. The feedback loop is direct and constitutive of the training. Artistic education is, in substantial part, an institutionalized practice for building the theory of the other that Tier 1 requires.

Language models are trained on the outputs of this process — on artifacts produced by people who have already internalized the orientation — but not on the feedback loops that produced the orientation itself. The resulting latent capacity is real but requires activation. The orientation sentence ("Think first about how the other person will experience this") is the activation key: it makes the criterion explicit at the point where the system begins generating, before domain conventions have organized the output around Tier 3 categories. The single sentence outperforms eleven domain principles precisely because it activates the capacity that is most latent and most consequential, rather than the capacity that is most explicit in the training signal.

The mechanism by which the orientation sentence produces its effect remains unestablished; three candidate accounts are consistent with the data. Under a *distribution-shift account*, the sentence narrows the generation distribution toward receiver-experience framings, suppressing domain-convention patterns. Under an *attractor-competition account*, orientation and domain-convention representations compete in early layers and the orientation sentence biases this competition. Under a *saliency-weighting account*, the sentence increases attention to experience-relevant task features, disproportionately represented in output. These accounts predict different intervention points and different transfer properties; distinguishing them requires interpretability methods beyond this study's scope. The computational mechanism of orientation is an open question.

6.3 The Symmetric Architecture

The same Tier 1 orientation that makes a good author makes a good reviewer, because both roles are ultimately evaluated by the same criterion: the quality of the downstream person's experience. A reviewer operating from Tier 3 — surfacing domain-correct observations without an orientation toward the solver or the author — produces feedback that is formally accurate and experientially irrelevant. The author cannot use it to improve the puzzle; the solver cannot use it to navigate the puzzle. C5's failure in the reviewing study (lens without philosophy = 21.7, below the 23.7 bare baseline) is the reviewing-side analog of the authoring-side L-T1 failure (experience-first removed = 12.3, categorical failure). Both failures occur for the same structural reason: Tier 3 without Tier 1 produces outputs organized around domain format rather than receiver experience.

This symmetry has implications beyond the two roles studied here. Any role in a creative production chain — editing, curating, directing, commissioning — that involves evaluating or shaping creative work for a downstream receiver should exhibit the same tier structure. The Tier 1 orientation, expressed appropriately for each role, is the universal entry point for quality in the chain.

This study has three limitations that bound the generalization. First, the empirical base is a single creative domain: puzzle-hunt design, which has unusual properties (fully verifiable solutions, established evaluation rubrics, a specific solver-experience criterion) that may not transfer directly to domains where quality criteria are more contested. Second, the authoring and reviewing agents are AI systems simulating human expert roles rather than human professionals responding to prompts; the transfer to human-in-the-loop creative systems requires separate validation. Third, this study examines only authoring and reviewing roles; other roles in the creative production chain — editing, directing, commissioning — may exhibit different tier weightings or different orderings within Tier 2.

7 Conclusion

Three empirical findings constitute the core contribution of this paper. First, creative quality in both authoring and reviewing roles is governed by a three-tier cognitive architecture: orientation (Tier 1, irreplaceable), operational discipline (Tier 2, load-bearing), and domain specificity (Tier 3, amplifying). These tiers are not additively independent — Tier 3 content without Tier 1 activation degrades performance below the no-context baseline, and Tier 2 components without Tier 1 direction operate on the wrong objective. Second, domain expertise is Tier 3 in both roles. It is necessary for achieving ceiling performance, substitutable across sources and formulations, and never sufficient without the orientation and discipline layers that determine what the domain knowledge is being used for. Third, the minimum sufficient creative prompt consists of three sentences: "Think first about how the other person will experience this. Remove anything unnecessary. Verify every claim." This three-sentence OLE formula — Orientation, Lens, Expertise — crosses the publishable quality threshold in the authoring study and is derivable from the reviewing-side tier structure independently.

The practical implication is direct. Any practitioner using an AI system for creative work can restructure their prompts to follow OLE order: orientation sentence first, discipline sentences second, domain content as context third — never as persona identity. The evidence base for this restructuring is quantitative and large. The orientation sentence alone accounts for a 15.6-point quality gap between the artist and puzzle designer personas on a 45-point scale. Its absence accounts for a 29.4-point catastrophic failure that produces categorical breakdown rather than degraded output. The minimum sufficient set (T1 + T4 + T5) reaches the publishable threshold at 35.7 out of 45 in the authoring study, and the equivalent structure (Tier 1 + three non-redundant Tier 2 principles) reaches the best-condition performance in the reviewing study. The OLE formula is not a heuristic extracted from subjective experience; it is a description of the empirically measured quality hierarchy across two roles in a controlled testbed.

Three directions for future work extend the tier model beyond puzzle-hunt design. First, cross-domain validation: does the OLE reconstruction curve replicate in business document generation, legal drafting, or interactive fiction? The tier model makes a strong prediction — orientation-first prompts will outperform domain-expertise-first prompts in any creative domain whose quality criterion is defined by receiver experience — and this prediction is falsifiable by any controlled authoring study that varies prompt structure rather than prompt content. The tier model's domain-independence is a structural prediction grounded in the universal role of creative receivers, not an empirically established finding; cross-domain replication is the required next step. Second, multi-agent creative systems: when a reviewing agent and an authoring agent both operate from the OLE structure, does the quality of their collaborative output exceed that of domain-labeled agents? The symmetric architecture of authoring and reviewing tiers predicts a compounding effect: an authoring agent oriented toward the solver and a reviewing agent oriented toward the solver and author produce a feedback loop in which Tier 1 orientation is reinforced at every exchange. Third, human validation: do human professionals who explicitly apply the OLE structure to their own creative briefs report quality improvements relative to their domain-expertise-first defaults? The tier model predicts yes in all three cases, and in all three cases the prediction is specific enough to be falsified. A tier model that fails to cross the threshold at the Tier 1 + Tier 2 boundary in a new domain, or fails to produce quality advantages in a multi-agent OLE system, would require revision. The value of the model is precisely its falsifiability — and the precision with which it specifies where and how domain expertise should enter a creative prompt.

References